

ELL783 Operating Systems

Assignment - 1

Bhawna Kumari (2018EE10454)

Preeti Saharan (2018EE10486)

Submitted: 27 February 2022

1 Introduction to the xv6 operating system

1.1 Installing and Testing xv6

QEMU and xv6 operating system were installed following the steps given in the assignment manual. A An extra package *qemu-system-x86* was installed by running the command *sudo apt install qemu-system-x86* to remove the error *Couldn't find a working QEMU executable* while running the command *make qemu*.

1.2 System calls

For adding a new system call following files have to be modified-

- *syscall.h* - In this file we map a system call name to a unique number like `#define SYS_ps 22`. This assigns number 22 to system call *ps* which lists the running tasks.
- *syscall.c* - This file contains the `syscalls[]` table where we give the pointer to the system call function. `[SYS_ps] sys_ps` - here `sys_ps` is the pointer. Another thing to add here is `extern int sys_ps(void);`
- *sysproc.c* - This file contains the implementation of the system call function and is called from `syscall()` function in *syscall.c*
- *usys.S* - It contains user level code. `SYSCALL(ps)`
- *user.h* - It contains the function signature of the system calls which the user can use. `int ps(void);`
- User program and Makefile - If there is a user program then to run it some changes have to be made in the Makefile as given in the assignment manual.
- *defs.h* - If any other *.c* file is modified with a function then the function's signature has to be included in the file *defs.h*. For example, `int ps()` is a helper function in *proc.c* file for *ps* system call, so it is included in *defs.h*.

1.2.1 Listing the running processes

System call name - *ps*

Function - Lists the currently running processes.

User program name - *process_list.c*

Command to run - *process_list*

Implementation - We will traverse through the whole page table and check if a process is in *RUNNING* state and if it is then we will print it.

Figure 1 shows the code snippet of *ps* and figure 2 shows the output of the user program.

```

680 int
681 ps(void)
682 {
683     struct proc *p;
684     acquire(&ptable.lock);
685     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
686         if(p->state == RUNNING)
687             cprintf("pid:%d name:%s\n", p->pid, p->name);
688     }
689     release(&ptable.lock);
690     return 22;
691 }

```

Figure 1: Code of *ps* in file *proc.c*

```

$ process_list
pid:4 name:process_list

```

Figure 2: Output of user program *process_list*

1.2.2 Printing the available memory

System call name - *memtop*

Function - tells the amount of memory available or free(in bytes) in the system

User program name - *available_mem.c*

Command to run - *available_mem*

Implementation - For this we have counted the number of free pages in the memory by traversing the *freelist* of the memory allocator and then multiplying the no. of free pages with the size of a page in bytes to get the total free available memory. This is done by creating a function *int availableMemory()* in the file *kalloc.c* which returns the available memory.

Figure 3 shows the code snippet of *memtop* and figure 4 shows the output of the user program.

```

97 int
98 availableMemory(void)
99 {
100     int count = 0;
101     struct run *r = kmem.freelist;
102     while(r){
103         count++;
104         r = r->next;
105     }
106     return count*PGSIZE;
107 }

```

```

$ available_mem
available memory: 232603648

```

Figure 4: Output of user program *available_mem*

Figure 3: Helper function for *memtop* in file *kalloc.c*

1.2.3 Context switching

System call name - *csinfo*

Function - To find the number of context switches of a process.

User program name - *context_switch.c*

Command to run - *context_switch*

Implementation - For this we have a new variable named *numSwitch* in the process struct in *proc.h* file to count the number of context switches. When a new process is created and allocated memory, then we initialize *numSwitch* = 0 in *allocproc()* function. Every time the scheduler runs a particular process, we increment its *numSwitch* by 1. So the system call just returns the process' *numSwitch* value.

Output - Since *cs1* and *cs2* are updated consecutively and there will not be in context switch of the process, so the first two values *cs1* and *cs2* will be same. Then we are calling sleep for 1 time unit so value of *cs3* will be one greater than *cs2* and again sleep is called so value of *cs4* will be one greater than *cs3*. Hence the output is like *cs1* = *cs2* = *x*(let) , *cs3*= *x*+1 and *cs4* = *x*+2.

Figure 5 shows the code snippet of *csinfo* and figure 6 shows the output of the user program.

1.3 Scheduling Policy

1.3.1 Current Scheduling Policy

- Code location - in *proc.c* file *scheduler* and *sched* functions
- Scheduling Policy - Round robin where each process runs until a timer interrupt.

```

105 int
106 sys_csinfo(void)
107 {
108     return myproc()->numSwitch;
109 }

```

Figure 5: Code for *csinfo* in file *sysproc.c*

```

$ context_switch
context switch counts = 8 8 9 10

```

Figure 6: Output of user program *context_switch*

- What happens when a process returns from an I/O operation? - for an I/O operation to happen the process goes into sleeping state and invokes *sched()* which executes *swtch* to switch context and save the current process's state and the control goes to *scheduler()* which will then loop over the process table to find a RUNNABLE process and start running it. After the IO operation is complete, wakeup system call changes the state of all the process sleeping to RUNNABLE so that they can be scheduled again.
- What happens when a new process is created and when/how often does the scheduling take place? - the *userinit()* function creates the first user process and after that all child processes are created by using *fork*. *allocproc()* is used to initialize the process parameters and memory. Since it is a round robin scheduler, each processor runs for about 100ms before a timer interrupt takes place.

1.3.2 First Come First Serve: (SCHEDFLAG= FCFS)

FCFS is a non-preemptive policy and the processes are run on the basis of their creation time. For this we have created a new variable *ctime* in process struct in *proc.h* file. When a process is created and allocated memory and its paramters are initialized in the function *allocproc()* then we define the *ctime* of the process as:

```
p->ctime = ticks;
```

where ticks is the counter for time used in xv6. Figure 7 shows the code for FCFS. In FCFS we traverse the process table and find the process with least creation time and run it. The ready times depend on the sequence in which processes are created and FCFS can lead to starvation.

1.3.3 Multi-level Queue Scheduling: (SCHEDFLAG= MLQ)

MLQ is a preemptive policy so for this we have used *QUANTA = 5* and the processes are run on the basis of their priority. For this we have created a new variable *priority* in process struct in *proc.h* file. When a process is created, it is assigned a priority 2. When a child process is created using *fork*, then the parent's priority is copied to the child also. In MLQ we traverse the process table and find the process with the highest priority and run it. To prevent starvation, round robin policy is followed in each priority queue by keeping track of the current process which is being executed and the next highest priority we need to execute . Figure 8 shows the code for MLQ. In this we can change the priority of the process using the system call *chpr* which takes the process pid and required priority as argument. Figure 9 shows the code for *chpr*.

1.3.4 Dynamic Multi-level Queue Scheduling: (SCHEDFLAG=DMLQ)

DMLQ is same as MLQ with the following modifications:

- We cannot use *chpr* to change priority manually.
- In file *exec.c* (line 101-105) which contains the code of *exec* system call, we check if the macro DMLQ is defined and if it is then we change the process' priority to 2.
- In the function *wakeup1()* (file *proc.c*, line 692-695) which is called when we return from sleeping mode, again we check if the macro DMLQ is defined and if it is then we change the process' priority to 1(highest priority).
- After every time quanta the priority is decreased by 1 if it is greater than 1. Figure 10 shows the snippet of code for this. In this code after every time quanta, for MLQ and DMLQ we yield the CPU and for DMLQ we also decrease the priority.

1.3.5 Computing statistics - ready, run, sleep and turnaround time

To find the ready time, run time and sleep time of a process we created three new variables *readytime*, *runtime*, *sleeptime* in process struct in *proc.h* and they are updated using the function *computeTime()* as shown in Figure 11. The function *computeTime()* traverses the process table and then depending on the state of the process it increases the ready (RUNNABLE state), run (RUNNING state) or sleep (SLEEPING state) time of that process. This function is run i.e. the time values are updated in every tick of the clock. To make these values available to the

```

430     #ifdef FCFS
431     acquire(&ptable.lock);
432
433     struct proc *proc_min_ctime = 0;
434
435     // Loop over the process table to find the process with minimum ctime
436     struct proc *current_process = ptable.proc;
437     while(current_process < &ptable.proc[NPROC]){
438         if(current_process->state == RUNNABLE){
439             if(proc_min_ctime == 0){
440                 proc_min_ctime = current_process;
441             }
442             else if(current_process->ctime < proc_min_ctime->ctime){
443                 proc_min_ctime = current_process;
444             }
445         }
446         current_process++;
447     }
448
449     if(proc_min_ctime){
450         // Switch to chosen process.
451         c->proc = proc_min_ctime;
452         switchvm(proc_min_ctime);
453         proc_min_ctime->state = RUNNING;
454
455         // Increase the no. of context switches.
456         proc_min_ctime->numSwitch = proc_min_ctime->numSwitch + 1;
457
458         swtch( &(c->scheduler),proc_min_ctime->context);
459         switchkvm();
460
461         // Process is done running for now.
462         // It should have changed its p->state before coming back.
463         c->proc = 0;
464     }
465
466     release(&ptable.lock);

```

Figure 7: Code for FCFS in file proc.c

user side, a new system call *waitstat* is created which is a slight modification of the wait system call and takes three integer pointers **readytime*, **runtime*, **sleeptime* as argument. In *waitstat* whenever a zombie process is found then during its clean up we update the values at these pointers so that we can retrieve them at the user end.

1.3.6 Testing the code

Performance comparison of all three policies

A user program *schedule_user* tests the three policies. In this program, we are yielding the CPU for short tasks based CPU bound processes so yield system call was implemented additionally. Table 1 shows the average statistics for all the scheduling policies and three types of tasks and figure 12 shows a sample output in terminal. To run it type command-

make clean qemu SCHEDFLAG=FCFS/MLQ/DMLQ
schedule_user (No of. processes to be created)

We have the following observations from the statistics:

- For IO processes, since we are running a loop of 1000 iterations and calling sleep system call, hence the sleep time for IO in all 4 policies is 1000. The sleep time for other cases is zero.
- For the three different types of process, the average ready time is in order: CPU <S-CPU <IO and this order is present in FCFS, MLQ and DMLQ. S-CPU tasks are yielded for 100 times so they spend more time in runnable state and for IO tasks sleep() system call is called 1000 times so when they return from sleep state then they go into runnable state hence we get this order.
- For the default scheduling policy, the average ready time is in the order: S-CPU <IO <CPU. This may be because of the frequent context switching due to timer interrupt which causes the CPU bound tasks to go into runnable state for more time.

Table 1: schedule_user - Average statistics for all policies for 10 processes

Policy	Process type	Ready time	Run Time	Sleep time	Turnaround time
DEFAULT	CPU	2005	4070	0	6075
DEFAULT	S-CPU	112	0	0	112
DEFAULT	IO	500	0	1000	1500
FCFS	CPU	1343	4319	0	5662
FCFS	S-CPU	2030	0	0	2030
FCFS	IO	4090	0	1000	5090
MLQ	CPU	1425	4497	0	5922
MLQ	S-CPU	2142	0	0	2142
MLQ	IO	3072	0	1000	4072
DMLQ	CPU	1727	5104	0	6831
DMLQ	S-CPU	3889	0	0	3889
DMLQ	IO	3960	0	1000	4960

- The average run time is non - zero for CPU bound processes in all the policies as we have added extra loops in the code to consume CPU time.
- For DMLQ as compared to other policies, the ready time of IO process is more or less same to that of S-CPU process because in DMLQ policy we are updating the priority of the process to highest after it returns from sleep.
- For CPU bound process we see that the performance is comparable for all 3 policies.

Performance of MLQ policy

A user program *MLQ_user_check* tests MLQ policy by creating n processes (taken as argument and n >5) and then changing their priority to either 1,2 or 3 depending on a condition. Table 2 shows the average statistics for each priority and overall average for 15 processes and figure 13 shows sample output in terminal. To run it type command-make clean qemu SCHEDFLAG=MLQ

MLQ_user_check (No of. processes to be created)

Based on the statistics, we have the following observations:

- The average ready time for the 3 priorities is in the order: 1 <2 <3. This is because we are executing all the processes with higher priority before moving to any lower priority process.
- Since there are no explicit sleep calls so the sleep time is zero.
- The run times are comparable for differen priorities.

Table 2: MLQ_user_check - Average statistics for all priorities for 15 processes

Priority	Ready time	Run Time	Sleep time	Turnaround time
1	98	43	0	141
2	101	42	0	144
3	258	52	0	311
Overall average	153	46	0	199

2 Introduction to Linux Kernel Modules

Header files

- linux/module.h - Need to include for every module
- linux/kernel.h - used it for printk
- linux/sched/signal.h - included it to make use of function "for_each_process()" to iterate over all the process.

Meta Information

MODULE_LICENSE("GPL") - GNU Public License

MODULE_AUTHOR("Preeti") - Author of the module
MODULE_DESCRIPTION("..") - Gives a short description about the module.

2.1 Kernel module writing

1. `int init_module(void)`

It is called whenever this module is loaded into kernel. It uses `printk` to print the message "Kernel Module Loaded". We used `KERN_INFO` log level (priority = 6). Returns 0 when successfully loaded. Figure 14 and Figure 15 shows corresponding code and output.

2. `cleanup_module(void)`

It is called whenever module is removed from kernel. Uses `printk` to print the message "Kernel Module Removed".

3. Commands

`make clean`

`make`

`sudo insmod KernelModuleWriting.ko` - inserted module

`dmesg -t—tail -1` - used `tail -1` to print recent 1 line to print the message written in `init_module`.

`sudo rmmod KernelModuleWriting` - removed module

2.2 Listing the running tasks

1. `int init_module(void)`

It is called whenever this module is loaded into kernel. Used `for_each_process()` to go through all the tasks, choose only those process which are runnable using the state of process. Created a function `get_task(state)` to obtain state (runnable, interruptible, or other). Then used `printk` to print task id, task command and task state. Figure 16 and Figure 17 shows corresponding code and output.

2. `cleanup_module(void)`

It is called whenever module is removed from kernel. Uses `printk` to print the message "ListingRunningTasks Kernel Module Removed".

5. Commands

`make clean`

`make`

`sudo insmod ListingRunningTasks.ko`

`dmesg`

`sudo rmmod ListingRunningTasks`

```

469     #ifdef MLQ
470     acquire(&ptable.lock);
471     struct proc *proc_high_priority = 0, *p, *new_sml_proc = sml_proc;
472     for(p = sml_proc; p < &ptable.proc[NPROC]; p++){
473         if(p->state == RUNNABLE){
474             if(proc_high_priority == 0){
475                 proc_high_priority = p;
476                 new_sml_proc = p;
477             }
478             if(p->priority < proc_high_priority->priority){
479                 proc_high_priority = p;
480                 new_sml_proc = p;
481             }
482         }
483     }
484     if(proc_high_priority){
485         c->proc = proc_high_priority;
486         switchvm(proc_high_priority);
487         proc_high_priority->state = RUNNING;
488         // No. of clock cycles run by the process in a quanta.
489         proc_high_priority->ticks_elapsed = 0;
490         // Increase the no. of context switches.
491         proc_high_priority->numSwitch = proc_high_priority->numSwitch + 1;
492         swtch(&(c->scheduler), proc_high_priority->context);
493         switchkvm();
494         // Process is done running for now.
495         // It should have changed its p->state before coming back.
496         c->proc = 0;
497
498         int next_max_priority = 3;
499         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
500             if(p->priority < next_max_priority){
501                 next_max_priority = p->priority;
502             }
503         }
504         if(next_max_priority == proc_high_priority->priority){
505             new_sml_proc++;
506             if(new_sml_proc < &ptable.proc[NPROC])
507                 sml_proc = new_sml_proc;
508             else
509                 sml_proc = ptable.proc;
510         }
511         else{
512             sml_proc = ptable.proc;
513         }
514     }
515     release(&ptable.lock);
516     #endif

```

Figure 8: Code for MLQ in file proc.c

```

782 // Changes priority of the process with given pid.
783 // Returns 1 after changing priority if a process is found with given pid.
784 // Returns -1 if no process found with given pid.
785 int
786 chpr(int pid, int priority)
787 {
788     struct proc *curr;
789     acquire(&ptable.lock);
790     for(curr = ptable.proc; curr < &ptable.proc[NPROC]; curr++){
791         if(curr->pid == pid){
792             curr->priority = priority;
793             release(&ptable.lock);
794             return 1;
795         }
796     }
797     release(&ptable.lock);
798     return -1;
799 }

```

Figure 9: Code for *chpr* in file *proc.c*

```

113 // Force process to give up CPU on clock tick.
114 // If interrupts were on while locks held, would need to check nlock.
115 #if defined(MLQ) || defined(DMLQ)
116 if(myproc() && myproc()->state == RUNNING &&
117     tf->trapno == T_IRQ0+IRQ_TIMER && (++myproc()->ticks_elapsed) == QUANTA){
118     #ifdef DMLQ
119         if(myproc()->priority > 1)
120             myproc()->priority -= 1;
121     #endif
122     yield();
123 }
124 #endif

```

Figure 10: Code run after every time QUANTA in file *trap.c*

```

710 void
711 computeTime(void)
712 {
713     struct proc *p;
714     acquire(&ptable.lock);
715     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
716         if(p->state == RUNNING)
717             p->runtime++;
718         else if(p->state == SLEEPING)
719             p->sleeptime++;
720         else if(p->state == RUNNABLE)
721             p->readytime++;
722     }
723     release(&ptable.lock);
724 }

```

Figure 11: Code for computing statistics in file *proc.c*


```

$ schedule_user 10
Printing statistics for each process -----
PID | Type | Ready time | Run time | Sleep time | Turnaround time
4   | S-CPU | 0           | 0         | 0           | 0
7   | S-CPU | 4           | 0         | 0           | 4
6   | CPU   | 0           | 4021      | 0           | 4021
9   | CPU   | 4           | 4016      | 0           | 4020
10  | S-CPU | 4020        | 1         | 0           | 4021
13  | S-CPU | 4022        | 0         | 0           | 4022
5   | IO    | 4019        | 0         | 1000        | 5019
8   | IO    | 4021        | 0         | 1000        | 5021
11  | IO    | 4023        | 0         | 1000        | 5023
12  | CPU   | 4020        | 4730      | 0           | 8750

Printing Average Statistics -----
Type | Avg Ready time | Avg Run time | Avg Sleep time | Avg turnaround time
CPU  | 1341           | 4255         | 0              | 5597
S-CPU | 2011           | 0            | 0              | 2011
IO    | 4021           | 0            | 1000           | 5021

```

Figure 12: Output for schedule_user with FCFS

```

$ MLQ_user_check 15
Printing statistics for each process -----
PID | Priority | Ready time | Run time | Sleep time | Turnaround time
15  | 1        | 0           | 56        | 0           | 56
16  | 2        | 9           | 54        | 0           | 63
18  | 1        | 61          | 40        | 0           | 101
19  | 2        | 62          | 40        | 0           | 102
21  | 1        | 102         | 40        | 0           | 142
22  | 2        | 105         | 39        | 0           | 144
24  | 1        | 144         | 40        | 0           | 184
25  | 2        | 146         | 40        | 0           | 186
27  | 1        | 186         | 39        | 0           | 225
28  | 2        | 187         | 39        | 0           | 226
17  | 3        | 227         | 46        | 0           | 273
20  | 3        | 229         | 46        | 0           | 275
23  | 3        | 267         | 41        | 0           | 308
26  | 3        | 269         | 40        | 0           | 309
29  | 3        | 301         | 90        | 0           | 391

Printing Average Statistics -----
Priority | Avg Ready time | Avg Run time | Avg Sleep time | Avg turnaround time
1        | 98             | 43           | 0              | 141
2        | 101            | 42           | 0              | 144
3        | 258            | 52           | 0              | 311

For all process with MLQ scheduling policy
Average ready time = 153
Average run time = 46
Average sleep time = 0
Average turnaround time = 199

```

Figure 13: Output for MLQ_user_check with 15 processes

```

#include <linux/module.h> /* need to include for every kernel module */
#include <linux/kernel.h> /* used it only for printk log levels */

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Preeti");
MODULE_DESCRIPTION("A simple module which prints KERNEL MODULE LOADED message when loaded and KERNEL MODULE REMOVED message when removed");

int init_module(void)
{
    printk(KERN_INFO "Kernel Module Loaded\n"); /*KERN_INFO priority = 6 , current console log level = 4*/
    return 0;

    /* 0 is returned when module is successfully loaded, in case of an error--> 0 will not be returned*/
}

void cleanup_module(void)
{
    printk(KERN_INFO "Kernel Module Removed\n");
}

```

Figure 14: Code for init and cleanup functions

```

[ 3739.580081] Kernel Module Loaded
[ 3772.309816] Kernel Module Removed

```

Figure 15: Output-printing module loaded and removed message

```

char * task_state(long state)
{
    switch (state) {
        case TASK_RUNNING:
            return "TASK_RUNNING";

        default:
            {
                return "TASK_NOT_RUNNING";
            }
    }
}

int init_module(void)
{
    struct task_struct *task; // Pointer to Task

    for_each_process(task)
    {
        if(strcmp(task_state(task->state), "TASK_RUNNING") == 0){ /* picking only the processes which are running*/
            printk( KERN_INFO "task_command: %s | state: %ld | process_id: %d\n", task->comm, task->state, task->pid );
        }
        /*printk( KERN_INFO "task_command: %s | state: %ld | process_id: %d \n", task->comm, task->state, task->pid);*/
    }

    return 0;

    /* 0 is returned when module is successfully loaded */
}

void cleanup_module(void)
{
    printk(KERN_INFO "ListingRunningTasks Kernel Module Removed\n");
}

```

Figure 16: Code for module of listing running tasks

```

[ 1171.884926] task_command: gnome-terminal- | state: 0 | process_id: 3397
[ 1171.884931] task_command: insmod | state: 0 | process_id: 4768

```

Figure 17: Output-printing running processes